Pre-analysis Plan: Betting Markets, Incentives, and Partisan Bias in Political Forecasting

May 5, 2026

1 Introduction

1.1 Research Motivation

In political forecasting, prediction markets can outperform raw opinion polls. Are betting markets able to improve over raw polling because they provide additional information, because they provide monetary incentives for accuracy, or because the act of betting changes how people treat their own beliefs? Our research aims to isolate whether incentives, information, or the act of betting itself changes how individuals form and report beliefs, thereby providing causal evidence to distinguish among these mechanisms.

1.2 Introduction to Betting Markets

Prediction markets, also known as betting markets or event futures, are exchanges where participants trade contracts whose payoff depends on the outcome of a future event. On platforms such as Kalshi or Polymarket, a typical contract pays \$1 if an event occurs and \$0 otherwise. The current contract price can therefore be interpreted as the market's collective probability estimate of that outcome. As of 2026, prediction markets such as Kalshi operate legally in the United States and host markets on politics, policy, economics, and other public events.

Prediction markets are understood in the literature as mechanisms for aggregating dispersed information into a single probabilistic forecast. They serve three related functions. First, they incentivize individuals to seek out information. Second, they encourage truthful revelation of beliefs through financial stakes. Third, they aggregate diverse opinions into a single, changing market price (Wolfers and Zitzewitz, 2004). Empirically, prediction markets have often outperformed traditional forecasting tools such as opinion polls, expert judgment, and consensus surveys, especially in political and economic markets. They also tend to incorporate new information quickly and are relatively resistant to manipulation, though they may still display behavioral biases (Arrow et al., 2008).

Experimental and theoretical work suggests that, under the right institutional conditions, market prices can converge toward rational expectations by aggregating privately held information among participants (Plott and Sunder, 1988). This implies that prediction markets can be viewed not only as betting platforms but also as information-processing systems. Even when no single individual has all relevant information, the market as a whole may generate accurate probability estimates.

1.3 Prediction Markets and Polls

In electoral forecasting, prediction markets can outperform raw opinion polls. The 2024 presidential election was a clear example of how prediction markets could outperform traditional polls that the United States had used for decades. In the months leading up to the election, many standard public opinion polls suggested a highly competitive race, with some polls even showing Kamala Harris leading in Iowa. However, prediction markets, particularly Kalshi, took a different view of this race. By late October 2024, Kalshi markets priced Donald Trump at a 61.8% chance of winning the election, compared to 38.2% for Harris. The total volume of trades in this prediction market exceeded $500 million by late October.

Berg, Nelson and Rietz (2008) show that, over the long run, prediction markets have often been more accurate than alternative forecasting methods. One intuitive explanation is that markets allow participants to put "skin in the game." Because participants can gain or lose money, they may have stronger incentives to set partisan beliefs aside and make forecasts based on evidence rather than expressive political preference.

Recent work on Polymarket similarly argues that markets can be more informative than polls. Market odds directly reflect candidates' perceived chances of winning because bettors put money behind their forecasts. As trading volume increases, odds can stabilize, especially in liquid markets where large trades do not move prices much. Markets also react more quickly than polls to important news events, such as debate performances or President Biden's withdrawal from the 2024 race (Tsang and Yang, 2026).

Diercks, Katz and Wright (2026) make a related point about macroeconomic event markets. Polls generally ask respondents whether they think an outcome will occur, while markets express a full probability for the event. For example, a market can imply a 30 percent chance that inflation is 3 percent and a 60 percent chance that it is 4 percent. This probabilistic format conveys both what is likely and how uncertain the outcome is. Polls usually do not reveal the same degree of confidence or risk.

These arguments suggest that financial incentives and probability-based reporting may explain part of the advantage of betting markets over polls. But understanding why polls fall short also requires looking at the human behavior that shapes survey responses, especially partisan bias.

1.4 Partisan Bias in Survey Responses

One possible reason opinion polls get things wrong is that respondents exhibit partisan bias in self-reported evaluations of facts and predictions about factual matters. Democrats and Republicans often evaluate the same facts differently, particularly when the question concerns the president, the national economy, or politically salient public events. After elections, supporters of the winning party often become more optimistic, while supporters of the losing party become more pessimistic, even when underlying economic conditions have not changed much.

Bartels (2002) uses panel data to show that partisanship shapes how people interpret political and economic information. Democrats and Republicans evaluate the same facts differently over time. This pattern helps explain why polls may capture partisan loyalty or worldview as much as objective assessments of reality.

Zilinsky and Bisbee (2026) analyze how party affiliation shapes economic perceptions and show that people report more favorable economic evaluations when their own party is in power and more negative evaluations when the opposing party is in power. This suggests that survey responses are filtered through partisan identity rather than based only on objective economic conditions.

Gerber and Huber (2009) show that partisan differences in economic perceptions can translate into real economic behavior. People become more optimistic when their preferred party is in power and more pessimistic when it is not, and these expectations are associated with behavior such as spending and investment. Partisan bias therefore affects not only survey responses but also economic action.

These patterns raise the key question for this study: if political identity distorts survey responses, can incentives or betting frames counteract that distortion?

1.5 Financial Incentives and Expressive Responding

Despite the many possible reasons that traditional polls fail, the success of prediction markets may come largely from financial incentives that motivate people to report more accurate beliefs. Current literature and experimental evidence show that incentives can reduce partisan bias in factual reporting.

Bullock and Lenz (2019) argue that partisan differences in factual beliefs are often smaller than previously assumed and frequently decline when respondents are incentivized for accuracy. Many respondents may have low confidence in their political knowledge and rely on partisan cues when uncertain. In ordinary surveys, respondents may engage in expressive responding, giving answers that signal allegiance to a political or social group rather than reporting their best estimate of the truth. This is sometimes called partisan cheerleading.

Across experimental studies, accuracy incentives tend to reduce some biased reporting. Rathje et al. (2023) use four preregistered survey experiments to test whether financial incentives improve respondents' ability to distinguish true from false political news headlines. They find that incentives improve accuracy and reduce partisan bias, especially by increasing belief in politically incongruent but factually true headlines.

Prior, Sood and Khanna (2015) conduct a survey experiment in which participants answer factual questions about the economy. Some participants receive monetary rewards for accuracy. They find that partisan bias declines when respondents are incentivized to be accurate. In ordinary surveys, respondents may treat questions as opinion opportunities; with accuracy incentives, they are more likely to treat them as factual questions.

The evidence is not uniformly one-sided. Berinsky (2018) runs experiments on politically controversial false beliefs about President Obama and finds limited evidence that incentives eliminate expressive responding. In one experiment, 42 percent of Republicans still reported that Obama was Muslim even when incentivized to say otherwise. This suggests that strong prior beliefs and cues may continue to influence responses, especially for controversial claims.

The present experiment builds on this literature by separating market information, accuracy incentives, and the psychological frame of betting.

2 Design

This study was deemed exempt by the University of Chicago Social & Behavioral Sciences Institutional Review Board (protocol IRB 26-0735). We will recruit 2,700 US-based respondents from the CloudResearch Connect survey platform. We oversample respondents ages 18–24 to better reflect the age profile of betting-market participants. Respondents are paid $1.50 for taking a six-minute Qualtrics survey.

2.1 Demographic Questions

We use demographic fields provided by CloudResearch Connect. The fields and response labels below match the Connect assignment export used for the study. Age is provided as a numeric field; the remaining fields are categorical labels in the Connect export.

1. **Age.** Numeric age in years. We include age in analysis as a continuous variable. Responses below 18 or above 100 will be flagged for review.

2. **Education.** For analysis, we include an indicator for bachelor's degree or higher. Respondents who select prefer not to say or have missing education are coded as zero on this indicator, and we include a separate indicator for prefer not to say or missing education.

3. **Employment Status.** For analysis, we include an indicator for respondents whose Connect employment status is Full-time or Business Owner. Respondents who select prefer not to say or have missing employment status are coded as zero on this indicator, and we include a separate indicator for prefer not to say or missing employment status.

4. **Ethnicity.** For analysis, we code a Hispanic indicator. Respondents with missing ethnicity are coded as zero on this indicator.

5. **Gender.** For analysis, we include an indicator for respondents coded as Man; Woman is the reference category. Respondents with missing gender are coded as zero on this indicator.

6. **Household Income.** We recode the Connect categories into a numeric variable by taking the midpoint of each observed income range. For the bottom category, less than $10,000, we code the midpoint as $5,000. For the top category, $250,000 or more, we code the midpoint as $275,000. Respondents who select prefer not to say or have missing income are assigned the sample median of the nonmissing income midpoint variable, and we include a separate indicator for prefer not to say or missing income.

7. **Race.** For analysis, we include indicators for White respondents and Black respondents. Respondents with missing race are coded as zero on both indicators. All other race categories are retained in the cleaned data but are not separately included in the main adjusted specifications.

We additionally collect the following study-specific questions.

8. **Party ID.** Generally speaking, do you think of yourself as a Republican, a Democrat, an Independent, or what? The first two answers are rotated randomly. Response options: Democrat, Republican, Independent, Other party, and No Preference.

   Respondents who select Democrat are asked whether they are strong Democrats or not very strong Democrats. Respondents who select Republican are asked whether they are strong Republicans or not very strong Republicans. Respondents who select Independent, Other party, No Preference, or skip the initial question are asked whether they think of themselves as closer to the Republican Party or to the Democratic Party, with Democrat and Republican rotated randomly and an additional option for neither.

   Party ID is coded on a seven-point scale from strong Democrat to strong Republican.

9. **Prior awareness of prediction markets.** Before today, had you ever heard of prediction markets such as Kalshi or Polymarket? Response options: yes, no, and not sure. Coding: yes is coded 1, no is coded 0, and not sure is coded missing.

10. **Prior use of prediction markets.** Have you ever placed money on a political prediction market or similar platform? Response options: yes, no, and unsure. Coding: yes is coded 1, no is coded 0, and unsure is coded missing.

2.2 Treatment and Measurement

We include the pre-test response in the demographic portion of the survey. The pre-test response looks like the control condition.

2.2.1 Instructions

At the pre-test, everyone sees:

> In this study, you will be asked to make predictions about several real-world events. For each event, please report the probability that you think the event will happen. A probability of 0% means you think the event definitely will not happen. A probability of 100% means you think the event definitely will happen. A probability of 50% means you think the event is equally likely to happen or not happen. Please answer as carefully and honestly as you can based on your own judgment.

2.2.2 Randomization

Respondents are randomized with equal probability across four conditions:

- Pure control
- Information only
- Accuracy incentive

- Betting frame

  At the post-test, everyone sees:

  > You will now again be asked to make predictions about the same real-world events as earlier in this survey.

  Everyone except the pure control group also sees:

  > You will also see a current market forecast for each event. This is the market's current estimate, not necessarily the true answer. You may use this information however you wish when making your own forecast. Your predictions on this survey will not impact the overall market forecast, are hypothetical, and are not affiliated with any market prediction platforms.

Information-only participants do not see bonus language, a bonus calculator, or any other monetary-incentive display.

Accuracy-incentive participants also see:

> You will earn a bonus based on the accuracy of your probability forecasts. For each event, your bonus contribution depends on how accurate your forecast is relative to the current market forecast. You will make forecasts for four events. Each forecast can contribute up to $0.25 to your total bonus, for a maximum total bonus of $1.00.
>
> Your final bonus will be calculated after the events are resolved.

On each event, the accuracy-incentive participants see a bonus calculator that shows how their forecast translates into a potential bonus contribution based on the current market probability and the accuracy of their forecast relative to the market.

Betting-frame participants see:

> You will earn a bonus based on the accuracy of your probability forecasts. For each event, your bonus contribution depends on how accurate your forecast is relative to the current market forecast.
>
> Imagine you have an experimental betting position. You may move your position above or below the market. Moving above the market is like betting YES. Moving below the market is like betting NO. The farther you move from the market, the larger your implied bet. You will make forecasts for four events. Each forecast can contribute up to $0.25 to your total bonus, for a maximum total bonus of $1.00.
>
> Your final bonus will be calculated after the events are resolved.

On each event, the betting-frame participants see the same bonus calculator as the accuracy-incentive participants.

2.2.3 Response

All respondents enter a probability forecast from 0 to 100 percent.

2.2.4 Attention Check

The pre-test forecast block includes one attention-check slider formatted like the event-forecast questions. The prompt instructs respondents who are reading the question to enter 64. We will code respondents who do not enter 64 as failing the attention check. Attention-check failure is a data-quality flag and is not an exclusion from the primary analysis.

2.2.5 Events

The survey will use four Kalshi markets. The event titles displayed in the survey are populated from the market-data endpoint at the time of the survey. The planned markets are:

Table 1: Planned Forecasting Events

Event	Market ID	Selection rule	Event text
1	`kxcpi-26apr`	Line: 0.5	Will CPI rise more than 0.5% in April 2026?
2	`kxpayrolls-26apr`	Line: 70,000	Will above 70,000 jobs be added in April 2026?
3	`controls-2026`	Contract: `CONTROLS-2026-R`	Will Republicans win the U.S. Senate in 2026?
4	`kxiceero-26mar`	Contract: `KXICEERO-26MAR-JUL01`	Will legislation that funds ICE's Enforcement and Removal Operations (ERO) become law before Jul 1, 2026?

2.3 Payoffs

Only respondents assigned to the accuracy-incentive and betting-frame conditions are eligible for an accuracy bonus. Respondents in the pure control and information-only conditions receive the base payment but do not receive a forecast-accuracy bonus and do not see a bonus calculator. No participant can lose money, and the maximum bonus is $1.

Let:

- $A_j = 1$ if event j happens and $A_j = 0$ otherwise.
- f_{ij} be participant i's forecast for event j.
- m_j be the current market-implied probability for event j.
- $S(f, A) = 1 - (A - f)^2$ be the quadratic accuracy score.

For bonus-eligible respondents, define the per-event score difference relative to the market benchmark as

$$\Delta S_{ij} = S(f_{ij}, A_j) - S(m_j, A_j),$$

which lies in $[-1, 1]$. The total bonus paid to respondent i is

$$\text{Bonus}_i = 0.5 + 0.5 \times \frac{1}{4} \sum_{j=1}^{4} \Delta S_{ij}.$$

This bounds the bonus in $[0, 1]$, with a \$0.50 midpoint and a \$1 cap. Forecasts that outperform the market benchmark earn higher bonuses, while forecasts that underperform the market benchmark earn lower bonuses. We use this rule instead of actual market trading because we want to elicit beliefs about the probability of the event, not gambling taste or loss aversion. The information-only condition is not bonus-eligible so that the design separates market information alone from market information with monetary incentives. The betting-frame condition is bonus-eligible so that respondents experience the betting frame in a setting with real accuracy incentives.

3 Analysis

3.1 Outcomes

3.1.1 Primary Outcome

The primary outcome is partisan divergence in post-treatment deviations from the market. For each respondent-event observation, define

$$D_{ij} = f_{ij} - m_{ij},$$

where f_{ij} is respondent i's post-treatment probability forecast for event j and m_{ij} is the market probability recorded at the post-test for that respondent and event. The market probability is displayed to respondents in the information-only, accuracy-incentive, and betting-frame conditions, but not to respondents in the pure control condition. Respondents enter forecasts on a 0–100 slider; in analysis, we divide these responses by 100 so that forecasts and market probabilities are both coded on the 0–1 probability scale. We then define the respondent-level mean market deviation

$$\overline{D}_i = \frac{1}{4} \sum_{j=1}^{4} D_{ij}.$$

The primary estimand is the treatment effect on the Republican-Democrat difference in $\overline{D}_i$. Let $R_i = 1$ for Republican or Republican-leaning respondents and $R_i = 0$ for Democrat or Democrat-leaning respondents. Respondents who are neither Republican-leaning nor Democrat-leaning are included in descriptive summaries but excluded from the primary partisan-divergence contrast.

For each treatment comparison, the unadjusted primary estimand is

$$\begin{aligned}&\left[E(\overline{D}_i \mid T_i = 1, R_i = 1) - E(\overline{D}_i \mid T_i = 1, R_i = 0)\right]\\&\quad - \left[E(\overline{D}_i \mid T_i = 0, R_i = 1) - E(\overline{D}_i \mid T_i = 0, R_i = 0)\right].\end{aligned}$$

Negative values mean the treatment reduces the Republican-Democrat gap in deviations from the market; positive values mean the treatment increases that gap. This estimand does not require classifying which event direction favors either party in advance. If both parties move in the same direction relative to the market, this will appear as common disagreement with the market rather than partisan divergence.

3.1.2 Baseline Measures

The baseline measures are defined per respondent i and averaged over events $j = 1, \ldots, 4$:

- **Mean pre-treatment market deviation:**

$$\overline{D}_i^{pre} = \frac{1}{4} \sum_j (f_{ij}^{pre} - m_{ij}).$$

We include this as the pre-test covariate in adjusted models for the primary outcome.

3.2 Analysis Sample and Missingness

The primary analysis sample will exclude respondents who do not consent and respondents missing the first post-test forecast. Respondents with attention-check failures, speed flags, straightlining flags, or age-review flags will be retained in the primary analysis unless otherwise stated; these flags will be reported descriptively. The attention-check flag equals 1 for respondents whose attention-check slider response is not 64 and equals 0 for respondents whose response is 64. Any respondent with missing or invalid treatment assignment will be excluded from treatment-contrast analyses.

We define post-treatment response as having a nonmissing first post-test forecast. If more than 5 percent of respondents in any randomized treatment group are missing post-treatment response, we will report inverse-probability weighted estimates as a missing-outcome sensitivity analysis. These weights will be based on group-specific models predicting the probability of post-treatment response conditional on pre-treatment covariates, including party identification, age, prior awareness, prior prediction-market participation, pre-treatment market deviation, and the simplified Connect demographic variables described above. The inverse-probability weights will be stabilized within treatment group.

For covariates, variables are coded as described above so that missing or prefer-not-to-say responses are represented explicitly where planned. If a respondent has no valid pre-test deviation values, we will impute `pre_dev_mean` to the sample median and include an indicator for imputed pre-test deviation in adjusted models.

3.3 Primary Analysis

All code for analysis is available in our GitHub repository: `https://github.com/UChicago-pol-methods/bet`.

The primary analysis will use the mean post-treatment market-deviation outcome, `dev_mean`, and will estimate how each treatment changes partisan divergence in this outcome. We will first report group means for the four randomized conditions and Republican-Democrat gaps in market deviations by condition, with standard errors and confidence intervals.

We will estimate three preregistered contrasts:

1. Information only minus control.
2. Any incentive minus control, where any incentive pools the accuracy-incentive and betting-frame conditions.
3. Betting frame minus accuracy incentive.

For each contrast, we will report an unadjusted difference in Republican-Democrat gaps and a Lin-style post-double-selection lasso adjusted estimate of the same difference-in-gaps contrast. The adjusted model is estimated within the relevant comparison sample, i.e., for the information-only contrast, the sample includes only control and information-only respondents and $T_i = 1$ for information-only respondents. For the any-incentive contrast, the sample includes control, accuracy-incentive, and betting-frame respondents and $T_i = 1$ for respondents assigned to either incentive condition. For the betting-frame contrast, the sample includes only accuracy-incentive and betting-frame respondents and $T_i = 1$ for betting-frame respondents.

The reported adjusted estimate will use a Lin-style cell structure by defining four cells: control Democrats, control Republicans, treated Democrats, and treated Republicans. The forced regressors in the final model are the non-baseline treatment-by-party cell indicators. Candidate adjustment terms are centered covariates and interactions between each non-baseline cell indicator and each centered covariate. We will select adjustment terms using the union of variables selected by lasso for the outcome and by lasso for each non-baseline cell indicator, then estimate the final model by OLS with the cell indicators forced in. The adjusted treatment effect on partisan divergence is the linear contrast

$$(\widehat{\mu}_{T,R} - \widehat{\mu}_{T,D}) - (\widehat{\mu}_{C,R} - \widehat{\mu}_{C,D}),$$

evaluated at the sample mean of the covariates. We report HC2 heteroskedasticity-robust standard errors.

As a descriptive check, we will report the average market probability for each event and raw Republican and Democrat mean deviations from the market in the control group for each event separately. This table will show whether partisan divergence is present in the uninformed control condition and whether both parties move in the same direction or opposite directions relative to the market.

As an additional pre-specified machine-learning estimator, we will estimate cross-fitted doubly robust scores using random-forest nuisance models for the conditional outcome under treatment, the conditional outcome under control, and the treatment propensity. For each respondent, the score is

$$\widehat{\psi}_i = \widehat{\mu}_1(X_i) - \widehat{\mu}_0(X_i) + \frac{T_i(Y_i - \widehat{\mu}_1(X_i))}{\widehat{e}(X_i)} - \frac{(1 - T_i)(Y_i - \widehat{\mu}_0(X_i))}{1 - \widehat{e}(X_i)}.$$

We will report the mean score among Republican-leaning respondents, the mean score among Democrat-leaning respondents, and the difference between those means. The standard error for the difference will be the empirical standard error of the corresponding subgroup-contrast influence score. The Lin-style double-lasso and random-forest doubly robust estimators are pre-specified analyses for the primary outcome rather than exploratory analyses.

4 Ethical Considerations

Informed consent. All survey respondents must know that they are participating in a study, what they will be doing in the survey, and how their responses will be used.

Debriefing. Since the survey is similar to a real prediction-market interface, all participants will receive a statement clarifying the hypothetical nature of the betting position and confirming that the study is unaffiliated with any real prediction-market platform:

> As stated at the beginning and throughout the survey, your predictions will not impact the overall market forecast, are hypothetical, and are not affiliated with any market prediction platforms. If you found questions related to prediction markets upsetting or are struggling with a gambling addiction we encourage you to reach out to 1-800-522-4700.

Financial compensation. No participant can earn a negative payoff. All participants receive a base payment. Respondents assigned to the accuracy-incentive or betting-frame condition can additionally earn a bonus based on the accuracy of their forecasts relative to the market benchmark, up to a maximum of \$1 and a minimum of \$0. Respondents in the pure control and information-only conditions do not receive an accuracy bonus.

Data privacy and anonymity. All responses, including political beliefs and preferences, will remain anonymous if study results are published.

A Planned Survey Text

This appendix gives the planned survey text for the betting-market data collection. Demographic fields are provided by CloudResearch Connect. The survey itself collects party identification, prior prediction-market experience, pre-test forecasts, randomized treatment instructions, post-test forecasts, and a debrief.

Respondents affirm informed consent and are paid $1.50 by CloudResearch for completing the study.

A.1 Party Identification

> Generally speaking, do you usually think of yourself as a Democrat, a Republican, an independent, or other?

Response options:

- Republican
- Democrat
- Independent
- Other party (please specify)
- No preference

Respondents who select Republican or Democrat are then asked:

> Would you call yourself a strong [selected party] or a not very strong [selected party]?

Response options:

- Strong
- Not Very Strong

Respondents who do not select Republican or Democrat are asked:

> Do you think of yourself as closer to the Republican Party or to the Democratic Party?

Response options:

- Closer to Republican Party
- Neither
- Closer to Democratic Party

A.2 Prediction-Market Experience

Respondents are asked:

> Before today, had you ever heard of prediction markets such as Kalshi or Polymarket?

Response options: Yes, No, Unsure.

Respondents are also asked:

> Have you ever placed money on a political prediction market or similar platform?

Response options: Yes, No, Unsure.

A.3 Pre-Test Forecasts

Before the pre-test forecast questions, respondents see:

> In this study, you will be asked to make predictions about several real-world events. For each event, please report the probability that you think the event will happen. A probability of 0% means you think the event definitely will not happen. A probability of 100% means you think the event definitely will happen. A probability of 50% means you think the event is equally likely to happen or not happen. Please answer as carefully and honestly as you can based on your own judgment.

Respondents then complete four slider questions, one per event:

> [Event title loaded dynamically]
>
> What do you think the true probability of this event is?

The slider label is "Probability." The slider runs from 0 to 100 and displays the selected value. Event titles are loaded dynamically from the selected prediction-market contracts.

The pre-test block also includes an attention-check slider:

> Event: The party participating in this survey (you) reads this question in its entirety.
>
> If you are reading this, answer 64 to the following question. What do you think the true probability of this event is?

A.4 Treatment Instructions

Respondents are randomized with equal probability to one of four betting-study treatment arms.

Control. Control-group respondents see:

> You will now again be asked to make predictions about the same real-world events as earlier in this survey.

Information only. Information-only respondents see:

> You will now again be asked to make predictions about the same real-world events as earlier in this survey.
>
> You will also see a current market forecast for each event. This is the market's current estimate, not necessarily the true answer. You may use this information however you wish when making your own forecast. Your predictions on this survey will not impact the overall market forecast, are hypothetical, and are not affiliated with any market prediction platforms.

Accuracy incentive. Accuracy-incentive respondents see:

> You will now again be asked to make predictions about the same real-world events as earlier in this survey.
>
> You will also see a current market forecast for each event. This is the market's current estimate, not necessarily the true answer. You may use this information however you wish when making your own forecast. Your predictions on this survey will not impact the overall market forecast, are hypothetical, and are not affiliated with any market prediction platforms.
>
> You will earn a bonus based on the accuracy of your probability forecasts. For each event, your bonus contribution depends on how accurate your forecast is relative to the current market forecast. You will make forecasts for four events. Each forecast can contribute up to $0.25 to your total bonus, for a maximum total bonus of $1.00.
>
> Your final bonus will be calculated after the events are resolved.

Betting frame. Betting-frame respondents see:

> You will now again be asked to make predictions about the same real-world events as earlier in this survey.
>
> You will also see a current market forecast for each event. This is the market's current estimate, not necessarily the true answer. You may use this information however you wish when making your own forecast. Your predictions on this survey will not impact the overall market forecast, are hypothetical, and are not affiliated with any market prediction platforms.
>
> You will earn a bonus based on the accuracy of your probability forecasts. For each event, your bonus contribution depends on how accurate your forecast is relative to the current market forecast.

> Imagine you have an experimental betting position. You may move your position above or below the market. Moving above the market is like betting YES. Moving below the market is like betting NO. The farther you move from the market, the larger your implied bet. You will make forecasts for four events. Each forecast can contribute up to $0.25 to your total bonus, for a maximum total bonus of $1.00.
>
> Your final bonus will be calculated after the events are resolved.

A.5 Post-Test Forecasts

All respondents complete four post-test forecast sliders, one for each event. Control-group respondents see the same screen format as in the pre-test:

> [Event title loaded dynamically]
>
> What do you think the true probability of this event is?

Respondents in the information-only, accuracy-incentive, and betting-frame groups see the same forecast question plus the current market probability:

> [Event title loaded dynamically]
>
> What do you think the true probability of this event is?
>
> The current market probability is [market probability]%.

Accuracy-incentive and betting-frame respondents additionally see:

> Move the slider to lock in your forecast and see this question's potential bonus contribution.
>
> If the event happens: This question contributes [bonus].
>
> If the event does not happen: This question contributes [bonus].

Information-only respondents do not see the bonus text or conditional bonus amounts. Control-group respondents do not see market information or bonus information. The market probability is fetched dynamically at survey runtime and saved as embedded data for all respondents who reach the post-test, including control-group respondents.

A.6 Debrief

At the end of the survey, respondents see:

> As stated at the beginning and throughout the survey, your predictions will not impact the overall market forecast, are hypothetical, and are not affiliated with any market prediction platforms. If you found questions related to prediction markets upsetting or are struggling with a gambling addiction we encourage you to reach out to 1-800-522-4700.

References

Arrow, Kenneth J., Robert Forsythe, Michael Gorham, Robert Hahn, Robin Hanson, John O. Ledyard, Saul Levmore, Robert Litan, Paul Milgrom, Forrest D. Nelson, George R. Neumann, Marco Ottaviani, Thomas C. Schelling, Robert J. Shiller, Vernon L. Smith, Erik Snowberg, Cass R. Sunstein, Paul C. Tetlock, Philip E. Tetlock, Hal R. Varian, Justin Wolfers and Eric Zitzewitz. 2008. "The Promise of Prediction Markets." *Science* 320(5878):877–878.

Bartels, Larry M. 2002. "Beyond the Running Tally: Partisan Bias in Political Perceptions." *Political Behavior* 24(2):117–150.

Berg, Joyce E., Forrest D. Nelson and Thomas A. Rietz. 2008. "Prediction Market Accuracy in the Long Run." *International Journal of Forecasting* 24(2):285–300.

Berinsky, Adam J. 2018. "Telling the Truth about Believing the Lies? Evidence for the Limited Prevalence of Expressive Survey Responding." *The Journal of Politics* 80(1):211–224.

Bullock, John G. and Gabriel Lenz. 2019. "Partisan Bias in Surveys." *Annual Review of Political Science* 22:325–342.

Diercks, Anthony M., Joshua D. Katz and Jonathan H. Wright. 2026. Kalshi and the Rise of Macro Markets. Working Paper 34702 National Bureau of Economic Research.

Gerber, Alan S. and Gregory A. Huber. 2009. "Partisanship and Economic Behavior: Do Partisan Differences in Economic Forecasts Predict Real Economic Behavior?" *American Political Science Review* 103(3):407–426.

Plott, Charles R. and Shyam Sunder. 1988. "Rational Expectations and the Aggregation of Diverse Information in Laboratory Security Markets." *Econometrica* 56(5):1085–1118.

Prior, Markus, Gaurav Sood and Kabir Khanna. 2015. "You Cannot Be Serious: The Impact of Accuracy Incentives on Partisan Bias in Reports of Economic Perceptions." *Quarterly Journal of Political Science* 10(4):489–518.

Rathje, Steve, Jon Roozenbeek, Jay J. Van Bavel and Sander van der Linden. 2023. "Accuracy and Social Motivations Shape Judgements of (Mis)Information." *Nature Human Behaviour* 7:892–903.

Tsang, Kwok Ping and Zhen Yang. 2026. "The Anatomy of Polymarket: Evidence from the 2024 Presidential Election.".
URL: *https://arxiv.org/abs/2603.03136*

Wolfers, Justin and Eric Zitzewitz. 2004. "Prediction Markets." *Journal of Economic Perspectives* 18(2):107–126.

Zilinsky, Jan and James Bisbee. 2026. "Economic Evaluations and Partisan Faultfinding: When Are Respondents Most Likely to Answer Survey Questions Honestly?" *Political Science Research and Methods* 14(2):506–522.